

CIFAR-10 Classification Project Documentation

Data Preprocessing and Feature Engineering

Preprocessing:

- Normalized all images using CIFAR-10 specific mean values [0.4914, 0.4822, 0.4465] and standard deviation values [0.247, 0.243, 0.261]
- Split the original training set into 80% training (40,000 images) and 20% validation (10,000 images)
- Configured a batch size of 256 for efficient training

Feature Engineering through Data Augmentation:

- Random horizontal flipping with 50% probability to introduce orientation variance
- Random rotation up to 20 degrees to simulate different camera angles
- Color jittering with brightness, contrast, and saturation adjustments of up to 10%
- Random sharpness adjustments with a factor of 2 (applied 20% of the time)
- Random erasing of small image regions (applied 75% of the time) to simulate occlusions
- Used standard normalization without augmentation for validation and test sets

Model Selection and Optimization Approach

Model Selection:

- Implemented a ResNet9 architecture, which provides a good balance between computational efficiency and performance
- Utilized residual connections to improve gradient flow and enable more effective training of deeper layers
- The architecture includes convolutional blocks with batch normalization and ReLU activations
- Designed a progression of channel depths (3→64→128→256→512) to extract increasingly complex features



Optimization Approach:

- Used Adam optimizer with an initial learning rate of 0.001
- Applied weight decay of 35e-5 as regularization to prevent overfitting

- Implemented gradient clipping at 0.0001 to stabilize training
- Employed a ReduceLROnPlateau scheduler to decrease learning rate by a factor of 0.3 when validation loss plateaued for 3 epochs
- Utilized early stopping with a patience of 5 epochs to prevent overfitting and save computational resources
- Evaluated model performance with multiple metrics: accuracy, precision, recall, and F1 score
- The final model is stored in *.pt* format for easy loading and inference during deployment. This approach ensures flexibility for various application scenarios.

Model Evaluation Overview

- The **loss and accuracy curves** indicate a steady improvement in model performance over 30 epochs, with both training and validation accuracy exceeding **90%**.
- The **precision, recall, and F1-score** show consistent growth, highlighting balanced predictions across different classes.
- The **confusion matrix** reveals strong classification accuracy, with most predictions aligning correctly, though some misclassifications occur in visually similar categories.



Overall, the model demonstrates **high accuracy and reliability**, effectively classifying CIFAR-10 images with minimal performance drops.

Report on Image Classification API Development

1. Data Preprocessing & Feature Engineering

- The input images are resized to 32x32 pixels to match the CIFAR-10 dataset format.
- Normalization is applied using the mean and standard deviation of the dataset:
 - Mean: [0.4914, 0.4822, 0.4465]
 - Standard Deviation: [0.247, 0.243, 0.261]
- Images are converted to tensors using `torchvision.transforms` for deep learning compatibility.

2. Model Selection & Optimization

- ResNet-9 architecture is used for classification due to its efficient performance on CIFAR-10.
- The model contains residual connections to enhance training stability.
- It is pre-trained on CIFAR-10 and fine-tuned using cross-entropy loss and an Adam optimizer.
- The model achieves high accuracy and is optimized for inference with `.eval()` mode.

3. Deployment Strategy

- **FastAPI** is used for efficient, lightweight API deployment.
- Create `.env` file with these `API_USERNAME="admin"` `API_PASSWORD="admin"` variables.
- Basic Authentication is implemented using `HTTPBasic`, with default credentials:
 - Username: admin
 - Password: admin
- The API is CORS-enabled to allow requests from any origin.
- The model is loaded at startup and runs on CUDA (GPU) if available for faster inference.
- Uvicorn is used as the ASGI server for high-performance API serving.

4. API Usage Guide

- Endpoint for API Status Check:
 - GET / → Returns `{"message": "Image Classification API is running"}`

- Endpoint for Image Prediction:
 - POST /predict/
 - Requires Basic Authentication (admin:admin).
 - Accepts an image file as input.
 - Returns: predicted_class with confidence